

MA 750 - Final Project

Nate Josephs

Matthew Wiens

Ben Draves

December 2, 2017

Abstract

One of the most fundamental tasks in Statistics is to understand the relationship between two random variables, X, Y , via an unspecified function $Y = f(X)$. Typically, $f(\cdot)$ is unknown and must be estimated from data relating X and Y . Estimating $f(\cdot)$ using maximum likelihood yields no meaningful solution when we consider *all* functions. Hence statisticians turn to estimating $f \in \mathcal{F}$ where \mathcal{F} is a function space with some structure that provides meaningful solutions to the problem at hand. In most cases however, these function spaces are fixed, with no regard to the sample from which we are trying to infer f . In order to utilize all information inherent in the data while still imposing structure on \mathcal{F} , *Sieve Estimation* allows \mathcal{F} to grow in complexity as n increases. Heuristically, as n increases, we attain a more robust understanding of f and should allow our modeling procedure to consider more complex forms of f . Sieve achieves this by introducing more complex functions to \mathcal{F} as n increases. Here, we consider the function space

$$\mathcal{F}_n = \left\{ g(x) : g(x) = \sum_{d=1}^{D(n)} \beta_d x^d \right\}$$

where $D(n) \rightarrow \infty$ as $n \rightarrow \infty$. We focus our efforts on estimating $D(n)$ as a function of the data. This report is organized as follows; in sections 1 and 2 we summarize some of the foundational results on Sieve estimation and introduce notation used throughout the report. In section 3 we introduce some theoretical applications and in section 4 we offer some methodologies of estimating $D(n)$. Lastly, in section 5 we analyze our methods via intensive simulation study and a real data application.

1 Introduction

2 Development of Sieves Estimation

2.1 Series Estimators

Sieve estimators are a general class of estimators that can applied to a number of non parametric problems; for example density estimation and regression. Within this framework,

a number of econometric models have been developed, such as for logistic regression, missing data, and measurement error. However, we will focus on a sieve estimator for non parametric regression as an example to explore in further detail.

Recall that a sieve estimator for the unknown mean function $g(x)$ in the regression is a sequence of models. The motivation of sieve estimators leads to additional constraints, namely that the models should be have finite dimension, and that the complexity (i.e. dimension) increases as n increases. Therefore the most natural choice for a sieve estimator and the choice we will focus on is a series estimator:

$$\hat{g}(x) = \sum_{d=0}^D \alpha_d \phi_d(x)$$

[Hansen] Such a formulation leads to traditional problems in linear models, so given the dimension of the model and a choice of basis function, finding the coefficients α has well known solutions. The constraints in the formulation are light; namely the choice of basis function must approximate functions in L^2 , among other natural constraints for the regression problem. For completeness, examples of non linear sieves are Neural Networks, or penalized sieve estimators. [Chen]

Given the sieve estimator defined as a series, there are two choices to be made: the dimension of the model, and the choice of basis function. D must be data dependent to scale the complexity of the estimator with the data. Therefore, for now we will assume there is an algorithm to choose D which produces a D_{opt} , for some sense of optimality. There are a number of choices for the basis functions, and we will highlight just a few possibilities here. We will discuss optimal choices of dimension in the following section.

2.2 Basis Functions

One popular choice of functional estimator is given by the series estimator with polynomial basis functions. Within the polynomial class, choices of coefficients uniquely determine the exact form for the estimate. For example the coefficient vector $\mathbf{c} = [0, 0 \dots, 1]$ would correspond to a basis of $\{x^k : k = 0, 1, 2, \dots\}$. The Hermite Polynomials would correspond to another choice of \mathbf{c} , which have desirable theoretical properties such as orthonormality. While a powerful class of estimators, polynomial basis functions are not a natural choice when $g(x)$ goes to zero at $\pm\infty$. Using a Fourier Series resolves this issue, which again highlights the importance of an appropriate choice of basis function in function estimator. With several applications in signal processing, the Fourier basis is proven to be a natural choice to approximate periodic functions. Lastly the Gaussian basis set, given by $\{\varphi^{(d)} : d = 0, 1, 2, \dots\}$ where $\varphi^{(d)}$ is d th derivative of the Normal density has several similar properties of the Fourier basis. Moreover, we see that these functions have a natural probabilistic interpretation (as do the Hermite Polynomials) which could be quite useful in inferential settings. The functional forms of these basis sets are given in the Table 2.2.

Where $\phi^{(d)}(x)$ is the d th derivative of the Normal density, which will be the Normal density with coefficients of the d th Hermite polynomial. Like the Fourier Series, the Gaussian

Basis Functions	Functional Form
Polynomials	$\sum_{i=0}^d c_d x^d$
Fourier	$a \cos(\pi dx) + b \sin(\pi dx)$
Gaussian	$\varphi^{(d)}(x)$

Table 1: Proposed series estimators with basis function ϕ for $\hat{g}(x)$

series estimator has the nice property that it goes to zero at $\pm\infty$.

Other common choices for the basis function are Splines and Wavelets. Different splines can be used as the sieve estimator, however the exact form and behavior depends on the choice of spline and the number of free parameters it has. This result can be shown by considering a constrained optimization problem over the squares of the D th derivatives of the class of potential sieve estimators. Also note that the choice of basis in the univariate case extends naturally to the multidimensional case, where the multivariate basis is constructed as a tensor product of the univariate basis. [Chen]

With the setup of a series estimator with a choice of dimension and basis function the sieve estimator has a number of similarities to the Kernel Density estimation problem. In both cases there are two choices to be made, one parameter that controls the bias variance tradeoff, and a second parameter that is related to underlying beliefs about the model. So the bandwidth in the Kernel estimation problem is analogous to the choice of dimension. Intuitively, as the dimension increases or the bandwidth decreases the estimator is more sensitive to local behavior and produces a rougher estimate. Similarly, the choice of the Kernel is analogous to the choice of basis function; the choice impacts the final model and exact statistical properties of the estimator, yet is less interesting because any reasonable choice of basis function or kernel function produces a similar estimate.

2.3 Analysis of Series Estimators

To begin our formal analysis of the series estimator suppose we have some function $f(x)$ that can be well approximated by a series estimator. That is we can write

$$y_i \equiv f(x_i) + e(x_i) = \sum_{d=0}^{\infty} \phi_d(x_i) \alpha_d + e(x_i)$$

where $\{\phi_d\}_{d=0}^D$ is a set of orthogonal basis functions and α_d are coefficients. Then the *sieves* estimator for this function is given by

$$\hat{f}(x_i) = \sum_{d=0}^D \phi_d(x_i) \alpha_d = \phi^T(x_i) \alpha$$

where $\phi^T(x_i) = (\phi_0(x_i), \phi_1(x_i), \dots, \phi_D(x_i))^T$ and $\alpha = (\alpha_0, \alpha_1, \dots, \alpha_D)$. Now organizing our matrices as follows

$$Y = \begin{bmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{bmatrix}, P_D = \begin{bmatrix} \phi_0(x_1) & \phi_1(x_1) & \dots & \phi_D(x_1) \\ \phi_0(x_2) & \phi_1(x_2) & \dots & \phi_D(x_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(x_n) & \phi_1(x_n) & \dots & \phi_D(x_n) \end{bmatrix}, \alpha = \begin{bmatrix} \alpha_0 \\ \alpha_1 \\ \vdots \\ \alpha_D \end{bmatrix}, e = \begin{bmatrix} e(x_1) \\ e(x_2) \\ \vdots \\ e(x_n) \end{bmatrix}$$

we can write our model as

$$Y = P_D \alpha + e$$

Seeing this a linear function in P_D gives rise to our OLS estimator of α

$$\hat{\alpha} = (P_D^T P_D)^{-1} P_D^T Y$$

From here we see that our estimates are given by

$$\hat{f}(x_i) = \phi^T(x_i) \hat{\alpha}$$

Due to the arsenal of statistical techniques for the linear regression model, we can extend this estimation procedure to account for correlated errors, weighted error structure, or even random effects. Now notice, that in the linear model framework these estimates would result in unbiased best estimates of the coefficients α , *given that we specify D correctly*. In the case where we do not specify the mean function correctly (i.e. misspecify D), our estimate become unbiased and the variance also increases. This again highlights in the importance of correct estimation of the dimension of the series estimator which we discuss in the proceeding section.

3 Estimating the Dimension D

3.1 Mean Integrated Squared Error

In this framework, we introduce two types of error. The first is given by the error of the approximation. This is, how well the series estimator matches the true underlying function. We will define this error by

$$r(x) = \phi^T(x) \alpha - f(x)$$

(*Notice that these are the true coefficient values*). The other source of error is of course the random variation around the true regression function. We will call this the residual error, $\epsilon(x_i)$. Then we see that from the original model $y_i = f(x_i) + e_i(x_i)$ that we can decompose the error as

$$e_i(x_i) = \epsilon(x_i) + r(x_i)$$

As we vary D we greatly reduce the variance in $r(x_i)$ as the more complex model will always account for more variance. But we need to be careful to ensure that the model is only

reducing $r(x_i)$ not $\epsilon(x_i)$. To try and find an analytical choice of D , we will find the Mean Integrated Squared Error (MISE). First consider the following expansion.

$$\begin{aligned}\hat{f}(x) - f(x) &= \phi^T(x)\hat{\alpha} - f(x) \\ &= \phi^T(x)\hat{\alpha} - \phi^T(x)\alpha + \phi^T(x)\alpha - f(x) \\ &= \phi^T(x)\hat{\alpha} - \phi^T(x)\alpha + r(x) \\ &= \phi(x)\phi^T(x)(\hat{\alpha} - \alpha) + r(x)\end{aligned}$$

Assume that the underlying X distribution is g . Then

$$\begin{aligned}MISE(D) &= \int (\hat{f}(x) - f(x))^2 g(x) dx \\ &= \int (\phi^T(x)(\hat{\alpha} - \alpha) + r(x))^2 g(x) dx \\ &= \int r(x)^2 g(x) dx + 2(\hat{\alpha} - \alpha) \int \phi^T(x)r(x)g(x) dx + (\hat{\alpha} - \alpha)^T \int \phi(x)\phi^T(x)g(x) dx (\hat{\alpha} - \alpha) \\ &= E(r^2(x)) + 2(\hat{\alpha} - \alpha)E(\phi^T(x)r(x)) + (\hat{\alpha} - \alpha)^T E(\phi(x)\phi^T(x))(\hat{\alpha} - \alpha)\end{aligned}$$

Now recall that the $r(x)$ was a projection error during our OLS estimate of α . Therefore, $r(x)$ and $\phi^T(x_i)$ exists in orthogonal spaces. Hence $E(\phi^T(x)r(x)) = 0$. Moro ever, since ϕ is a collection of orthogonal functions the off diagonal elements of $E(\phi(x)\phi(x)^T)$ are all zero. This gives

$$\begin{aligned}MISE(D) &= E(r^2(x)) + \text{tr} \left[(\hat{\alpha} - \alpha)^T E(\phi(x)\phi(x)^T)(\hat{\alpha} - \alpha) \right] \\ &= E(r^2(x)) + \text{tr} \left[E(\phi(x)\phi(x)^T) E((\hat{\alpha} - \alpha)(\hat{\alpha} - \alpha)^T) \right]\end{aligned}$$

Now according to Hansen 2012 with $E(e(x)^2|x) = \sigma_x^2$, $\mathcal{Q} = E(\phi(x)\phi^T(x))$ and $\Omega = E(\phi(x)\phi^T(x)\sigma_x^2)$ due to the asymptotic behavior of this quantity we have

$$MISE(D) \simeq E(r^2(x)) + \frac{1}{n} \text{tr} (\mathcal{Q}^{-1}\Omega)$$

Now, if we assume homoskedasticity, we see that $\Omega = \sigma^2 \mathcal{Q}$. So plugging into the equation above we see that

$$MISE * (D) \simeq E(r^2(x)) + \frac{1}{n} \text{tr} (\sigma^2 \mathcal{Q}^{-1} \mathcal{Q}) = E(r^2) + \frac{\sigma^2}{n} \dim(\mathcal{Q}) = E(r^2(x)) + \frac{\sigma^2 D}{n}$$

Hansen 2012 then show that

$$MISE * (D) = MISE(D)(1 + o(1))$$

so $MISE*$ is close to $MISE$. But notice, even in this simple case we still require estimation of $r(x)$, which is directly related to the estimation of the true regression mean function

$f(x)$ which is what we started to estimate in the first place. Therefore, while an excellent theoretical tool, using $MISE^*$ is not feasible. But as we will see later on we have

$$MISE \approx PRESS \approx E(CV)$$

Therefore, we will use cross validation instead to choose D but will have a theoretical justification of our approach.

Develop PSE and connect to MISE (N)

3.2 Prediction Squared Error

Instead of MISE, one may be interested in calculating the Predicted Square Error (PSE) in order to find the optimal dimension with respect to PSE. If x^* is a new value from $X \sim f(x)$, then our prediction of Y given $X = x^*$ under the sieve estimator is $\hat{y}^* = \hat{f}(x^*)$. We then define PSE as the expectation of the squared error between Y^* , the actual value of the regression line at $X = x^*$, and \hat{y}^* . Observe that

$$\begin{aligned} PSE(\hat{f}(x^*)) &= E[(Y^* - \hat{y}^*)^2] \\ &= E\left[\left(f(x^*) + e^* - \hat{f}(x^*)\right)^2\right] \\ &= E\left[\left(e^* + (f(x^*) - \hat{f}(x^*))\right)^2\right] \\ &= E[e^{*2}] + 2E\left[e^*(f(x^*) - \hat{f}(x^*))\right] + E\left[\left(f(x^*) - \hat{f}(x^*)\right)^2\right] \\ &= E[(e^* - E[e^*])^2] + 2E[e^*]E\left[(f(x^*) - \hat{f}(x^*))\right] + E\left[\left(f(x^*) - \hat{f}(x^*)\right)^2\right] \\ &= Var(e^*) + 0 + \int E[(f(x) - \hat{f}(x))^2]g(x)dx \\ &= Var(e^*) + MISE(\hat{f}(x)) \end{aligned}$$

Hence, the optimal dimension for our sieve estimator with respect to PSE will be the same as the optimal dimension with respect to MISE, since minimizing PSE is equivalent to minimizing MISE. Note that in our derivation, we use the fact that the errors are zero mean and e^* is independent of the estimator.

3.3 Choosing D via Cross Validation

If we define $\tilde{e} = Y^* - \hat{y}^*$, then $PSE(\hat{f}(x^*)) = E[\tilde{e}^2]$, which we may interpret as the expectation of a single leave-one-out (LOO) squared prediction error, where our estimator

is fit on X_1, \dots, X_n and validated against $X = x^*$. This motivates us to consider LOO prediction errors for each $i = 1, \dots, n$, which will reveal a data-driven process for choosing an optimal dimension D .

For each i , define $\tilde{e}_i = y_i - \hat{y}_{(i)}$ where $\hat{y}_{(i)}$ is fit on $X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_n$. Then $PSE(\hat{f}_{(i)}(x)) = E[\tilde{e}_i^2]$ and we define the cross-validation (CV) criterion as

$$CV(\hat{f}) = \frac{1}{n} \sum_{i=1}^n \tilde{e}_i^2$$

By the linearity of expectation, we have that $E[CV(\hat{f})] = PSE(\hat{f}(x^*))$, which we will utilize as our data-driven procedure for choosing D_{opt} . This estimate for $MISE$ has asymptotic results that overcome the infeasibility of finding D_{opt} analytically.

4 Simulation Study

5 Real Data Application

6 Conclusion